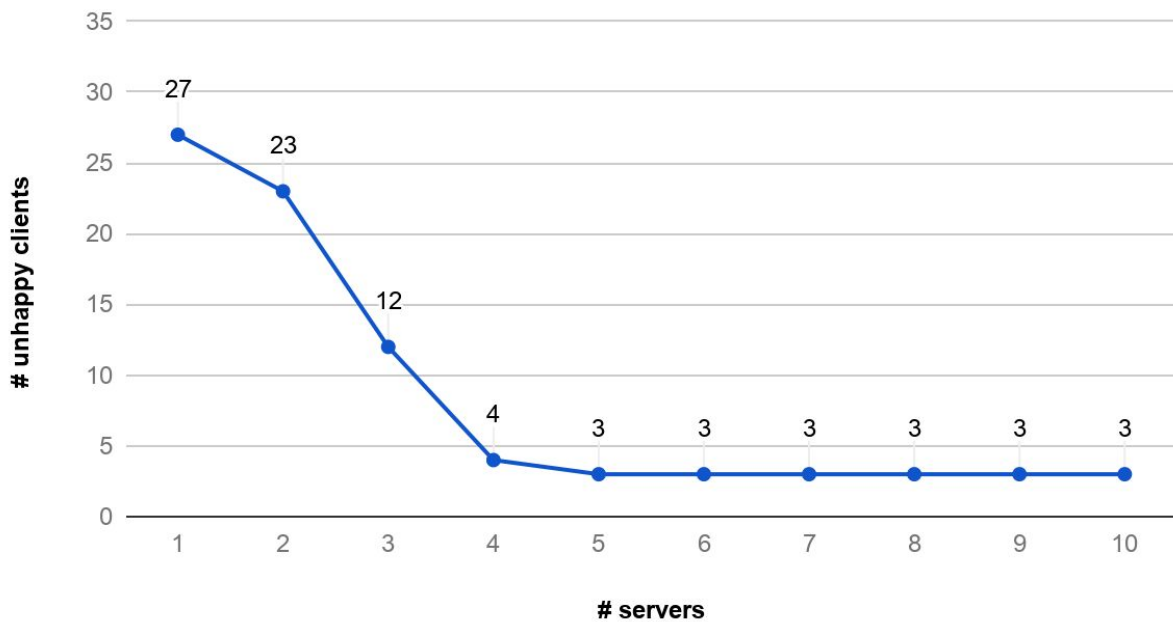


Benchmark Plot



servers	1	2	3	4	5	6	7	8	9	10
unhappy clients	27	23	12	4	3	3	3	3	3	3

Testing environment: `rand_spec = c-1000-42 hour = 1`

The plot above has a bend point at 4 servers. After adding into 3 servers in total, the number of unhappy clients (timeouts) drops significantly and when it comes to 4 servers, the performance considering VM time and timeout is basically the best. After 4 servers, the curve just flattens out.

So for this load, 3 servers can basically handle normal or non-busy situations and 4 servers can deal with busy hours. That is also the design for my checkpoint 1 server code.